

## ← Module 4 Quiz

Quiz, 10 questions

1  
point

1. Which of the following is an example of clustering?

- ☒ Separate the data into distinct groups by similarity
- ☐ Accumulate data into groups based on labels
- ☐ Creating a new representation of the data with fewer features
- ☐ Compress elongated clouds of data into more spherical representations

1  
point

2. Which of the following are advantages to using decision trees over other models? (Select all that apply)

- ☐ Decision trees can learn complex statistical models using a variety of kernel functions
- ☒ Trees often require less preprocessing of data
- ☐ Trees are naturally resistant to overfitting
- ☒ Trees are easy to interpret and visualize

1  
point

3. What is the main reason that each tree of a random forest only looks at a random subset of the features when building each node?

- ☐ To reduce the computational complexity associated with training each of the trees needed for the random forest.
- ☐ To learn which features are not strong predictors
- ☐ To increase interpretability of the model
- ☒ To improve generalization by reducing correlation among the trees and making the model more robust to bias.

1  
point

4. Which of the following supervised machine learning methods are greatly affected by feature scaling? (Select all that apply)

- ☒ KNN

- ☐ Decision Trees
- ☒ Neural Networks
- ☐ Naive Bayes
- ☒ Support Vector Machines

1  
point


5. Select which of the following statements are true.

- ☐ For a fitted model that doesn't take up a lot of memory, **KNN** would be a better choice than **logistic regression**.
- ☐ For having an audience interpret the fitted model, a **support vector machine** would be a better choice than a **decision tree**.
- ☒ For a model that won't overfit a training set, **Naive Bayes** would be a better choice than a **decision tree**.
- ☐ For predicting future sales of a clothing line, **Linear regression** would be a better choice than a **decision tree regressor**.

C: wrong  
D: wrong

1  
point

6. Match each of the prediction probabilities decision boundaries visualized below with the model that created them.

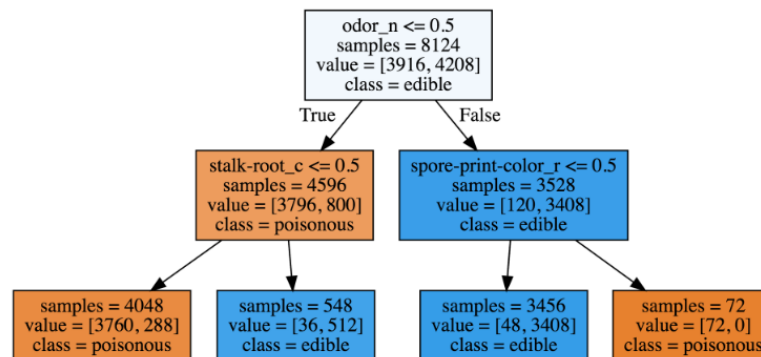


- ☐ 1. KNN ( $k=1$ )
- ☐ 2. Decision Tree
- ☐ 3. Neural Network
- ☐ 1. Neural Network
- ☐ 2. Decision Tree
- ☐ 3. KNN ( $k=1$ )
- ☒ 1. Neural Network
- ☐ 2. KNN ( $k=1$ )
- ☐ 3. Decision Tree

- ☐ 1. KNN (k=1)
- ☐ 2. Neural Network
- ☐ 3. Decision Tree

1  
point

7. A decision tree of depth 2 is visualized below. Using the `value` attribute of each leaf, find the accuracy score for the tree of depth 2 and the accuracy score for a tree of depth 1.



What is the improvement in accuracy between the model of depth 1 and the model of depth 2? (i.e. accuracy2 - accuracy1)

0.068

$$(3760+512+3408+72)/8124 - (3796+3408)/8124$$

1  
point

8. For the autograded assignment in this module, you will create a classifier to predict whether a given blight ticket will be paid on time (See the module 4 assignment notebook for a more detailed description). Which of the following features should be removed from the training of the model to prevent data leakage? (Select all that apply)

- ☒ compliance\_detail - More information on why each ticket was marked compliant or non-compliant
- ☐ ticket\_issued\_date - Date and time the ticket was issued
- ☐ graffiti\_status - Flag for graffiti violations
- ☐ agency\_name - Agency that issued the ticket
- ☒ collection\_status - Flag for payments in collections

1  
point

9. Which of the following might be good ways to help prevent a data leakage situation?

- ☐ If time is a factor, remove any data related to the event of interest that doesn't take place prior to the event.
- ☐ Ensure that data is preprocessed outside of any cross validation folds.
- ☒ Remove variables that a model in production wouldn't have access to
- ☐ Sanity check the model with an unseen validation set


A,C: Wrong  
C,D: wrong  
C: Wrong



1  
point

10. Given the neural network below, find the correct outputs for the given values of  $x_1$  and  $x_2$ .

The neurons that are shaded have an activation threshold, e.g. the neuron with  $>1?$  will be activated and output 1 if the input is greater than 1 and will output 0 otherwise.



x1	x2	output
0	0	1
0	1	0
1	0	0
1	1	1



x1	x2	output
0	0	0
0	1	1
1	0	1
1	1	1



x1	x2	output
0	0	0
0	1	1
1	0	1
1	1	0



x1	x2	output
0	0	0
0	1	0
1	0	0
1	1	1

- 
- ☒ I, **JINLONG Ban**, understand that submitting work that isn't my own may result in permanent failure of this course or deactivation of my Coursera account.

[Learn more about Coursera's Honor Code](#)

Submit Quiz

